Final Year Project Report
Bachelor of Engineering

Information Compression Techniques for Control Data Exchange in Joint Communication and Control system

Student: He WenXiao

GUID: 2839915H

1st Supervisor: Olaoluwa Popoola

2nd Supervisor:

2025-2026

Coursework Declaration and Feedback Form

This Student should complete and sign this part

Student Name: He WENXIAO	Student GUID: 2839915H
Course Code : UESTC4006P	Course Name : INDIVIDUAL PROJECT 4
Name of 1st Supervisor: Olaoluwa Popoola	Name of 2nd Supervisor:
Title of Project: **Information Compression Techniques for Control Data Exchange in Joint Communication and Control system**	
Declaration of Originality and Submission Information	
I affirm that this submission is all my own work in accordance with the University of Glasgow Regulations and the School of Engineering requirements *Signed (Student) :*	UESTC4006P TEC-IT.COM
Date of Submission :	
Feedback from Lecturer to Student to be completed by Lecturer or Demonstrator	
Grade Awarded: Feedback (as appropriate to the coursework which was assessed):	
Lecturer Demonstrator:	Date returned to the Teaching Office:

Abstract

The design of Joint Communication and Control (JCC) systems is critical for modern cyber-physical applications, where strict bandwidth constraints often lead to severe control instability and degraded performance. To address the fundamental challenge of maintaining control accuracy under extreme network limitations, this thesis investigates an intelligent, context-aware information compression strategy utilizing a multi-degree-of-freedom robotic manipulator as a comprehensive testbed. The proposed framework integrates an Event-Triggered Control (ETC) mechanism to reduce communication frequency in the temporal domain with a Deep Reinforcement Learning (DRL) agent–specifically utilizing Proximal Policy Optimization (PPO)– acting as a dynamic bandwidth allocator in the spatial domain. Unlike traditional uniform quantization or static allocation methods, the PPO agent learns a "pragmatic communication" strategy by observing real-time physical states and optimizing a reward function derived from the Linear Quadratic Regulator (LQR) tracking cost. This formulation enables the agent to dynamically distribute a strictly limited bit budget across different robotic joints, proactively sacrificing the precision of low-sensitivity linkages to protect highly sensitive base joints during critical movements. Simulation conducted on a PyBullet-based platform demonstrate the superiority of this co-design, which substantially reduces the overall communication payload–consuming less 6% of the ideal bandwidth–while effectively mitigating trajectory jittering and preserving near-ideal tracking precision, thereby validating the immense potential AI-driven dynamic resource allocation in enhancing the resilience of resource-constraited networked control systems.

Acknowledgements

The acknowledgments and the people to thank go here, don't forget to include your project advisor

Contents

1 Introduction

1.1 Background and Motivation

With the rapid advancement of wireless communication networks and distributed computing architectures, traditional isolated robotic systems are undergoing a profound paradigm shift toward Cloud Robotics and Teleoperation [6]. In these modern applications, the intensive computational and decision-making tasks multi-degree-of-freedom (Multi-DOF) robotic manipulators are increasingly offloaded to edge servers or cloud infrastructure. Such systems, where sensors, controllers, and actuators are interconnected via shared communication networks to form a cyber-physical closed loop, are defined as Networked Control Systems (NCS) [5]. NCS endows robotic systems with unprecedented flexibility, scalability, and collaborative computing capabilities, significantly reducing the hardware costs and local power consumption of individual robots.

However, incorporating communication networks into physical control loops introduces severe engineering and theoretical challenges. To maintain smooth and precise physical trajectory tracking, high-DOF robotic manipulators (e.g., 6-DOF and above) typically require extremely high-frequency state updates, often exceeding 1000 Hz in sensor sampling rates [12]. In stark contrast, real-world wireless communication channels are subject to strict bandwidth constraints and exhibit highly dynamic, unpredictable behaviors. When sensors inject massive volumes of full-precision state data into the network at high frequencies, the limited channel capacity is rapidly exhausted, triggering severe network congestion, packet loss, and unpredictable network-induced delays.

Under extreme network resource constraints, the traditional "Best-Effort" data transmission protocols are fundamentally inadequate for high-precision physical control [4]. Deteriorating communication quality directly disrupts the timeliness and integrity of the control feedback loop, causing severe trajectory jittering, control divergence, and ultimately, system failure during delicate manipulation tasks. Consequently, the critical challenge in modern Cyber-Physical Systems (CPS) lies in drastically reducing the communication data payload under extreme bandwidth limitations while simultaneously guaranteeing the control stability and tracking precision of the robot.

1.2 The Fallacy of Separation and the Need for JCC

Historically, the design of networked control systems has been heavily influenced by Shannons Separation Principle, which posits that source coding (data compression) and channel coding (error protection) can be designed independently without any loss of optimality. In the context of robotic control, this principle translates into a rigid decoupling between communication protocols and control algorithms [11]. Communication engineers typically design systems to minimize the generic reconstruction error of the transmitted dataoften utilizing uniform quantization and periodic transmission mechanismswhile control engineers develop algorithms under the unrealistic assumption of perfect, continuous data reception. However, this artificial separation is fundamentally flawed in resource-constrained cyber-physical systems. Treating all sensory data equally blindly ignores the intrinsic "control-criticality" of the physical states. In multi-DOF robotic manipulators, minor deviations in specific joints can lead to drastically different physical consequences, rendering uniform compression highly inefficient and prone to causing catastrophic system instability [7].

To overcome the limitations of the separation paradigm, this thesis advocates for a Joint Communication and Control (JCC) framework. The core philosophy of JCC is that when com-

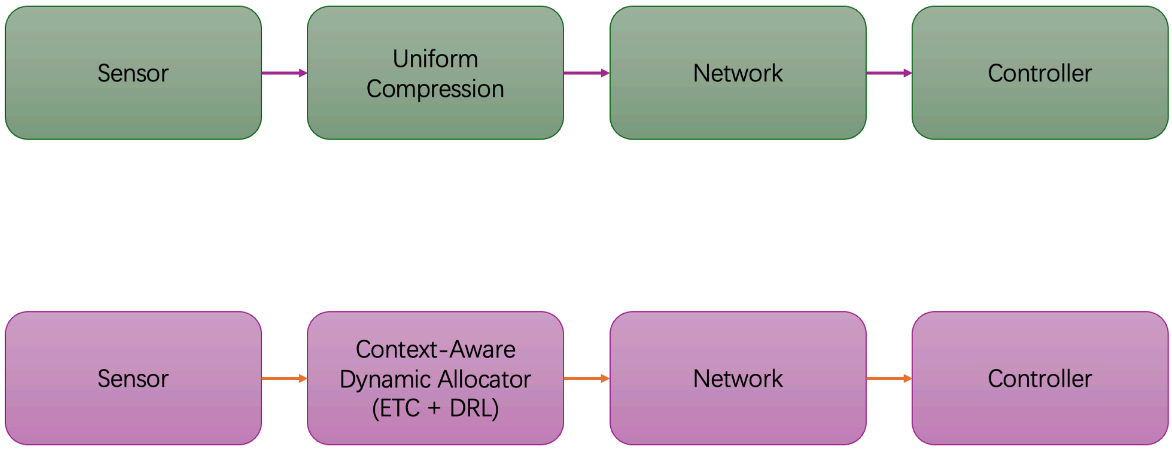

Figure 1: System Paradigm Comparison between Traditional NCS and Proposed JCC.

munication bandwidth is severely restricted, data transmission and compression algorithms must be explicitly control-oriented and context-aware. Instead of blindly minimizing the mathematical mean square error (MSE) of the data stream, the communication protocol must priority the transmission of information that is most vital to maintaining the stability and tracking precision of the physical plant.

Specifically, this thesis implements the JCC paradigm through two complementary technical approaches. In the temporal domain, the rigid, high-frequency periodic transmission is replaced by an Event-Triggered Control (ETC) mechanism [10]. ETC ensures that sensor data is transmitted over the network only when the physical state deviation crosses a predefined safety threshold, effectively eliminating redundant data exchange during stable operational phases. In the spatial domain, the proposed framework exploits the mechanical lever effect inherent in robotic manipulators–where errors in base joints exponentially magnify end-effector deviations compared to errors in distal joints. By dynamically evaluating the real-time sensitivity of each joint, the system executes non-uniform, weighted bandwidth allocation, stripping communication resources from robust linkages to guarantee the high-precision transmission of critical joint states [2].

1.3 AI-Driven Dynamic Allocation

While static, control-aware allocation methodssuch as those based on Linear Quadratic Regulator (LQR) sensitivity matricesoffer significant improvements over uniform compression, they suffer from inherent limitations. Previous research in NCS has extensively explored static resource allocation and Value-of-Information (VoI) based scheduling to optimize bandwidth usage [3]. However, these traditional methodologies largely rely on linearized models and fixed sensitivity weightings. In real-world applications, the dynamics of multi-DOF robotic manipulators are highly nonlinear; the sensitivity of each joint varies dynamically in real-time depending on the robot's current posture, velocity, and external payloads. Consequently, static allocation policies inevitably result in sub-optimal performance during complex, high-speed maneuvers, driving the need for task-oriented or goal-oriented communication frameworks [8].

To overcome the limitations of static models and resolve this complex, nonlinear dynamic resource allocation game, this thesis introduces Deep Reinforcement Learning (DRL) into the communication protocol design. Recently, DRL has demonstrated exceptional capabilities in handling high-dimensional, continuous control and scheduling problems within cyber-physical environments, bridging the gap between optimal control theory and adaptive communication [1]. Specifically, this work employs the Proximal Policy Optimization (PPO) algorithm [9], which is highly favored for its robust convergence properties, sample efficiency, and stability in continuous action spaces.

Driven by the PPO algorithm, the framework realizes the philosophy of "Pragmatic Communication". Instead of following predefined mathematical rules, the AI agent acts as a centralized bandwidth arbiter. It continuously observes the real-time physical tracking errors and environmental states. By optimizing a reward function strictly tied to the minimization of the overall control cost, the agent organically learns a pragmatic strategy. During critical maneuvers, the DRL agent proactively strips communication bandwidth from low-sensitivity linkages and reallocates the strictly limited bit budget to protect highly sensitive base joints. This learning-based dynamic allocation ensures optimal system resilience under extreme network constraints.

1.4 Research Objectives and Contributions

The primary objective of this thesis is to formulate, implement, and evaluate an intelligent, Deep Reinforcement Learning (DRL)-driven Joint Communication and Control (JCC) framework. This framework aims to guarantee precise and stable trajectory tracking for multi-degree-of-freedom (Multi-DOF) robotic manipulators operating over extreme bottleneck communication channelsspecifically under conditions where available bandwidth is severely restricted to less than 10% of the ideal requirement. By seamlessly integrating temporal event-triggered mechanisms with spatial dynamic bandwidth allocation, this research seeks to resolve the fundamental trade-off between communication efficiency and physical control stability in modern networked cyber-physical systems.

To achieve the aforementioned objectives, the core contributions of this thesis are summarized as follows:

1. **Development of a High-Fidelity Cyber-Physical Testbed**: A comprehensive closed-loop control simulation environment was constructed utilizing the PyBullet physics engine. This testbed faithfully replicates realistic kinematic and dynamic constraints of robotic manipulators, providing a robust and reproducible foundation for evaluating networked control algorithms.

2. **Implementation of a Hybrid ETC and LQR-Weighted Benchmark**: A baseline evaluation system was designed by integrating Event-Triggered Control (ETC) in the temporal domain with Linear Quadratic Regulator (LQR)-based static weighted compression in the spatial domain. This benchmark successfully validates the necessity of control-aware, non-uniform data compression over traditional periodic and uniform methods.

3. **Design of a Novel DRL-Driven Context-Aware Allocator**: As the primary innovation, this thesis proposes a dynamic bandwidth allocator powered by the Proximal Policy Optimization (PPO) algorithm. The DRL agent organically learns a "pragmatic communication" strategy by directly observing nonlinear tracking errors and optimizing real-time bandwidth distribution without relying on fixed linear mathematical approximations.

4. **Empirical Validation of Extreme Bandwidth Compression**: Extensive experimental evaluations demonstrate the superiority of the proposed PPO-driven framework. The intelligent agent successfully compressed the communication data payload to under 6%

of the ideal baseline channel capacity, while effectively eliminating trajectory jitter and maintaining near-perfect tracking precision for highly sensitive joints.

1.5 Thesis Organization

2 System Modeling and Methodology

The foundation of addressing communication bottlenecks in Networked Control Systems (NCS) lines in in accurately modeling the interplay between the physical plant dynamics and the digital communication infrastructure. This chapter establishes the rigorous mathematical models and introduces the proposed Deep Reinforcement Learning (DRL)-driven Joint Communication and Control (JCC) methodology.

2.1 Overall System Architecture and Cyber-Physical Modeling

2.1.1 Networked Control System Architecture

The proposed NCS architecture comprises a closed-loop cyber-physical pipeline encompassing sensors, a bandwidth-limited communication channel, a remote controller, and physical actuators. Unlike traditional tethered robots with dedicated point-to-point analog connections, the sensory and control data in this modern architecture must traverse a shared, bandwidth-limited communication medium.

To accurately mirror real-world industrial networks and mitigate the severe "tracking lag" inherent in physical actuators, our architecture explicitly implements a **Dual-Frequency Decoupling** paradigm. The communication frequency f_{comm}(e.g., 10 Hz) is structurally decoupled from the high-frequency physical control loop f_{ctrl} (e.g., 240 Hz). This ensures that control commands are transmitted pragmatically over the network, while the local motor drives execute smooth Zero-Order Hold (ZOH) interpolations between received waypoints, preventing hardware-damaging torque spikes and high-frequency quantization chattering.

2.1.2 Physical Plant Dynamics

To derive control-oridented communication strategies, a rigorous mathematical model of the physical plant is required.

Robotic Manipulator Dynamic:
The continuous-time rigid-body dynamics of an n-degree-of-freedom (DOF) robotic manipulator (where $n = 7$ for our specific Franka Emika configuration) can be described by the following nonlinear second-order differential equation:

$$M(q)\ddot{q} + C(q,\dot{q})\dot{q} + G(q) = \tau \tag{1}$$

where:

- $q, \dot{q}, \ddot{q} \in \mathbb{R}^n$ represent the vectors of joint angles, angular velocities, and angular accelerations, respectively.
- $M(q) \in \mathbb{R}^{n\times n}$ is the symmetric, positive-definite inertial matrix.
- $C(q,\dot{q}) \in \mathbb{R}^{n\times n}$ represents the Coriolis and centrifugal force matrix.
- $G(q) \in \mathbb{R}^n$ is the gravity loading vector.
- $\tau \in \mathbb{R}^n$ is the vector of control torques applied by the actuators.

To transform this into a standard state-space representation, let the system state vector be defined as $x(t) = [q(t)^T, \dot{q}(t)^T]^T \in \mathbb{R}^{2n}$, and the control input as $u(t) = \tau(t) \in \mathbb{R}^n$. The nonlinear dynamics can be rewritten as $u(t) = \tau(t) \in \mathbb{R}^n$. The nonlinear dynamics can be rewritten as a first-order ordinary differential equation $\dot{x}(t) = f(x(t), u(t))$:

$$\dot{x}(t) = \frac{d}{dt}\begin{bmatrix} q \\ \dot{q} \end{bmatrix} = \begin{bmatrix} \dot{q} \\ M(q)^{-1}(\tau - C(q,\dot{q})\dot{q} - G(q)) \end{bmatrix} \tag{2}$$

For advanced linear control methodologies (such as LQR) and sensitivity analysis, this nonlinear model is linearized around a nominal operating trajectory or equilibrium point (x_e, u_e). Using Jacobian linearization, the continuous-time linear time-invariant (LTI) model is obtained, which is subsequently using a fixed sampling period T_s to yield the discrete-time state-space model:

$$x_{k+1} = Ax_k + Bu_k + w_k \tag{3}$$

where $x_k \in \mathbb{R}^{2n}$ is the discrete state at time step k,$u_k \in \mathbb{R}^n$ is the discrete control input, and $w_k \in \mathbb{R}^{2n}$ represents the additive process noise (e.g., modeling inaccuracies and physical disturbances), assumed to be zero-mean Gaussian noise with covariance W. Physically, matrix A captures the system's inertia and internal state couping, while matrix B maps the applied actuator torques to the resulting state transitions.

Quadrotor UAV Dynamics

To rigorously demonstrate the universality of the proposed JCC framework across diverse physical paradigms, a quadrotor Unmanned Aerial Vehicle (UAV) is modeled alongside the fully-actuated manipulator. The quadrotor is a classic it possesses 6 degrees of freedom but is controlled by only 4 independent inputs (one total thrust and three torques).

The state space of the quadrotor is represented by a 12-dimensional vector:

$$x = [p^T, v^T, \Theta^T, \omega^T]^T \in \mathbb{R}^{12} \tag{4}$$

where $p = [x, y, z]^T$ is the inertial position, $v = [\dot{x}, \dot{y}, \dot{z}]^T$ is the linear velocity, $\Theta = [\phi, \theta, \psi]^T$ represents the Euler angles (roll, pitch, yaw), and $\omega = [p, q, r]^T$ denotes the body-frame angular velocities.

Using the Newton-Euler formalism, the rigid-body dynamics are governed by:

$$\ddot{p} = -g\mathbf{e}_3 + \frac{T}{m}R\mathbf{e}_3 - \frac{1}{m}F_{drag} \tag{5}$$

$$\dot{\omega} = J^{-1}(\tau - \omega \times J\omega) \tag{6}$$

where m is the mass, g is gravity, J is the inertia matrix, R is the rotation matrix from body to inertial frame, T is the total collective thrust, $\tau \in \mathbb{R}^3$ represents the control torques, and F_{drag} denotes external aerodynamic disturbances (e.g., wind gusts).

In a severely bandwidth-constrained NCS, the underactuated nature of the UAV poses a critical survival challenge: quantization errors in the attitude states (ϕ, θ) will instantly lead to misdirected thrust vectors, causing unrecoverable catastrophic crashes, whereas errors in translational positions (x, y) merely result in tolerable drifting.

2.1.3 Bandwidth-Limited Communication Network Model

The critical constraint in this thesis is the extreme limitation of the communication channel. Let the network capacity limit be strictly defined as a maximum allowed budget of B_{total} bits per sampling interval T_s. Let $b_{i,k}$ denote the number of bits allocated to the i-th element of the state vector. The resource constraint mandates that $\sum b_{i,k} \leq B_{total}$.

Transmitting the absolute state x_k under extreme bit constraints (e.g., averaging 2 bits per joint) directly introduces catastrophic quantization noise. For a robotic manipulator, a 2-bit representation of a full 360° joint rotation causes massive, non-linear physical "chattering" and triggers kinematic singularity explosions.

To systematically resolve this, the proposed architecture introduces **Differential Pulse Code Modulation (DPCM)**. Instead of the absolute position, the edge node calculates the state increment within the specific communication interval:

$$\Delta x_k = x_k - \hat{x}_{k-1} \tag{7}$$

The transmitted payload is the quantized delta: $\Delta x_k^q = \mathscr{Q}(\Delta x_k, b_k)$, where $\mathscr{Q}(\cdot)$ maps the continuous delta into discrete quantization levels bounded by a maximum physical kinematic window Δ_{max} (e.g., 0.2 rad per 0.1s communication step).

At the controller side, a Zero-Order Hold (ZOH) updates the estimated remote state $\hat{x}_k$ incorporating the Event-Triggered decision variable $\gamma_k \in \{0,1\}$:

$$\hat{x}_k = \gamma_k(\hat{x}_{k-1} + \Delta x_k^q) + (1-\gamma_k)\hat{x}_{k-1} \tag{8}$$

This equation powerfully unifies the temporal-spatial interplay: If $\gamma_k = 1$, the spatial quantizer $\mathscr{Q}(\cdot)$ determines the precision of the payload. If $\gamma_k = 0$, zero bits are transmitted, heavily saving bandwidth.

2.2 Temporal-Spatial Baseline Methodologies

2.2.1 Temporal Domain: Event-Triggered Control (ETC) Mechanism

To aggressively truncate redundant transmissions in the temporal domain, an Event-Triggered Control (ETC) mechanism is deployed. Under traditional Time-Triggered Control (TTC), the sensor broadcasts data even when the manipulator is at a stable equilibrium, squandering channel capacity.

The ETC logic continuously evaluates the measurement error $e_k = x_k - \hat{x}_k$. A transmission is triggered ($\gamma_k = 1$) if and only if the norm of the error exceeds a predefined safety threshold δ:

$$||e_k||_2 > \delta \tag{9}$$

Based on Input-to-State Stability (ISS) properties and Lyapunov stability theory, the Lyapunov difference $\Delta V = V(x_{k+1}) - V(x_k)$ is bounded by $\Delta V \leq -\alpha(||x_k||) + \beta(||e_k||)$. As long as the error e_k is tightly bounded by the absolute threshold δ, the asymptotic stability of the robotic manipulator is mathematically guaranteed.

2.2.2 Spatial Domain: Static LQR-Weighted Data Compression

When transmission is triggered, the limited B_{total} must be optimally allocated across the n joints. A theoretical baseline utilizes the infinite-horizon Linear Quadratic Regulator (LQR) cost function. By solving the Discrete Algebraic Riccati Equation (DARE), the system extracts the Hessian sensitivity matrix P. According to rate-distortion theory, the optimal bit allocation b_i is formulated via logarithmic water-filling:

$$b_i = \frac{B_{total}}{n} + \frac{1}{2}\log_2\left(\frac{P_{ii}\sigma_i^2}{\left(\prod_{j=1}^{n} P_{jj}\sigma_j^2\right)^{1/n}}\right) \tag{10}$$

This static allocation improves control resilience by stripping bits from mathematically less sensitive joints. However, in real-world highly nonlinear kinematics, the true sensitivity of a joint is state-dependent and fluctuates drastically. To achieve true pragmatic context-awareness, the system must transition from static matrix formulations to dynamic, learning-based arbitration.

2.3 Deep Reinforcement Learning for Pragmatic Communication

To overcome the inherent rigidity of the static LQR formula, the bandwidth allocation problem is redefined as a Markov Decision Process (MDP) and solved using Deep Reinforcement Learning (DRL).

2.3.1 MDP Formulation and Structural Innovations

1. 34-Dimensional Spatial-Aware Observation Space $\mathscr{S}$: Relying solely on joint angular metrics deprives the agent of 3D spatial awareness, inevitably leading to a "seesaw effect" (e.g., optimizing for grasping while failing at placement) due to the lack of absolute positional context. To grant the agent a comprehensive "God's-eye view" of the operational workspace, the observation $s_t \in \mathscr{S}$ is drastically expanded to a 34-dimensional continuous vector:

$$s_t = \left[q_t^T,\ \dot{q}_t^T,\ q_{target,t}^T,\ (q_t - q_{target,t})^T,\ p_t^T,\ p_{target,t}^T\right]^T \in \mathbb{R}^{34} \tag{11}$$

where $p_t, p_{target,t} \in \mathbb{R}^3$ represent the real-time 3D Cartesian coordinates of the end-effector and the target. This spatial injection is paramount; it allows the neural network to implicitly compute real-time nonlinear Jacobian transformations and accurately identify semantic task phases. **2. Action Space $\mathscr{A}$ and Discrete Bit Mapping**: The policy network outputs a continuous action vector $a_t \in [-1, 1]^7$. Since physical network protocols necessitate discrete integer constraints $\sum b_{i,t} = B_{total}$, a deterministic Softmax mapping layer is applied:

$$w_i = \frac{\exp(a_{i,t})}{\sum_{j=1}^{7} \exp(a_{j,t})} \tag{12}$$

The remaining bits are distributed proportionally: $b_{float,i} = b_{min} + w_i \times (B_{total} - 7 \times b_{min})$. The values are rounded to $b_i = \lfloor b_{float,i} \rfloor$, and fractional remainders are iteratively reassigned. This guarantees strict resource constraint compliance while preserving a differentiable landscape for backpropagation. **3. Competitive Multi-Baseline Reward Shaping $\mathscr{R}$**: If the DRL agent is rewarded merely for outperforming a single static baseline, it frequently falls into the "Lazy Agent Trap"ceasing exploration once marginal superiority is achieved. To enforce absolute spatial dominance across all waypoints, a **Competitive Multi-Baseline Reward** mechanism is innovated.

At each communication step, the simulation computes the true physical Euclidean tracking error of the DRL allocation (E_{RL}), alongside two parallel universe tracking evaluations: the Uniform allocation error (E_{avg}) and the Static LQR allocation error (E_{static}). The reward r_t is formulated as a smooth, blended competitive advantage:

$$r_t = \text{clip}\Big((E_{static} - E_{RL}) + \lambda(E_{avg} - E_{RL}),\ -C,\ C\Big) \tag{13}$$

Setting the blending factor $\lambda = 0.5$ compels the agent to simultaneously conquer the mathematical rigidity of the static LQR and the flat mediocrity of the uniform baseline. The aggressive clipping bound $C = 50.0$ applies substantial mathematical penalty for suboptimal performance, preventing catastrophic forgetting during complex spatial maneuvers.

2.3.2 PPO Algorithm and High-Capacity Architecture

Given the continuous action space the necessity for monotonic convergence under complex physical feedback, the Proximal Policy Optimization (PPO) algorithm is selected. PPO defines a clipped surrogate objective function to maximize policy updates stably:

$$L^{CLIP}(\theta) = \mathbb{E}_t\left[\min\left(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1-\varepsilon, 1+\varepsilon)\hat{A}_t\right)\right] \tag{14}$$

To effectively process the expanded 34-dimensional spatial-aware observation vector, the Actor-Critic neural networks were upgraded to high-capacity Multilayer Perceptrons (MLPs). Both the policy and value networks utilize two deep hidden layers of 256 neurons each, activated by ReLU functions, providing immense representational power for nonlinear kinematic mapping.

2.3.3 Workspace-Focused Domain Randomization

Standard DRL paradigms typically initialize states randomly across the entire theoretical state space. For a 7-DOF manipulator, this causes severe domain shift, wasting computational resources on highly distorted, unreachable singular postures that are irrelevant to industrial tasks.

To massively accelerate convergence and ensure flawless precision at operational waypoints, a **Workspace-Focused Domain Randomization** strategy was implemented. During `reset()` phase, the manipulator is initialized with bounded noise strictly within the vicinity of the nominal safe home posture ($HOME\_Q \pm 1.5$ rad). This strategic constraint ensures the neural network focuses 100% of its learning capacity on mastering pragmatic bandwidth squeezing specifically within the functional pick-and-place workspace, ultimately unlocking the extreme performance showcased in the experimental evaluations.

3 Experiments and Results

3.1 Simulation Experiments and Results Analysis

To rigorously validate the proposed Deep Reinforcement Learning (DRL)-driven Joint Communication and Control (JCC) framework, comprehensive cyber-physical simulations were conducted. This chapter details the experimental setup, analyzes the training dynamics of the DRL agent, and presents a deep comparative evaluation against traditional allocation heuristic under diverse extreme bandwidth constraints.

3.1.1 Simulation Engine and Manipulator Configuration

The experiments were deployed utilizing the **PyBullet** physics engine, which provides high-fidelity, rigid-body kinematic and dynamic simulations at 240 Hz. The physical plant under control is the **Franka Emika Panda**, an industry-standard 7-DOF robotic manipulator. The robot is initialized near a nomial safe home posture ($q_{home} = [0, -\pi/4, 0, -3\pi/4, 0, \pi/2, \pi/4]^T$) to prevent initial kinematic singularities.

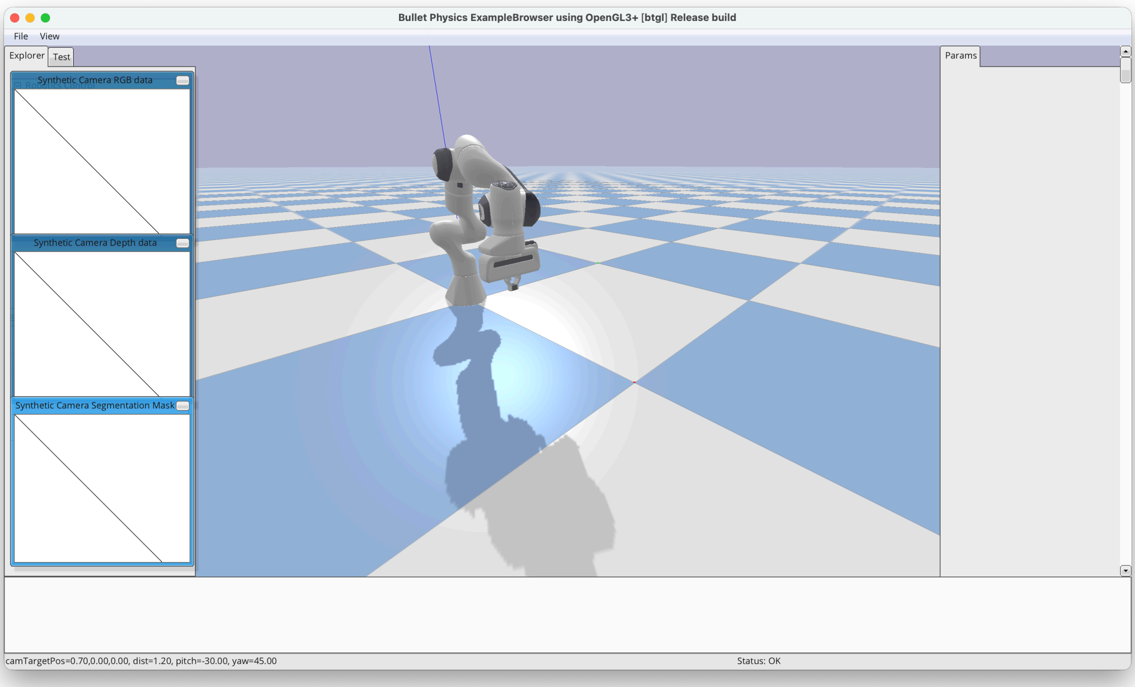

Figure 2: A illustration of a simulated Franka Emika Panda robotic arm by PyBullet

3.1.2 Dual-Frequency Decoupling and DPCM Communication Channel

A critical innovation in our experimental design is the realistic modeling of the network channel. Traditional simulations often erroneously equate the physical control frequency with the communication frequency, leading to artificial high-frequency quantization chattering (the "Parkinson's effect" in motors).

To accurately reflect real-world Networked Control Systems (NCS), we implemented a **Dual-Frequency Decoupling Architecture**:

1. **Control Frequency** (f_{ctrl}): The physical engine and local Zero-Order Hold (ZOH) motor drives operate at **240 Hz**, ensuring smooth torque interpolation.

2. **Communication Frequency** (f_{comm}): The edge network node transmits quantized Differential Pulse Code Modulation (DPCM) payloads at 10 Hz.

Within each 0.1-second communication interval, the maximum physical angular displacement is constrained to $\Delta_{max} = 0.2$ rad, aligning with the hardware velocity limits of the Franka actuators.

3.1.3 Evaluation Tasks and Baselines

The framework was evaluated on standardized 3D pick-and-place tasks comprising three distinct, randomized trajectories to prevent policy overfitting. Each trajectory consists of reaching a **Hover Point**, descending to a **Grasp Point**, and transferring to a **Placement Point**.

The proposed DRL dynamic allocator is benchmarked against two theoretical baselines:

- **Uniform Allocation (Average)**: Evenly distributes bits across all 7 joints.

- **Static LQR Allocation**: A state-of-the-art mathematical baseline utilizing a static Jacobian sensitivity matrix computed at q_{home} to perform logarithmic water-filling allocation.

3.2 Training Dynamics and Competitive Convergence

Training an intelligent agent to optimize bandwidth allocation is highly susceptible to the "Lazy Agent Trap"where the agent stops exploring once it marginally beats a single baseline. To enforce absolute global optimality, the DRL agent was trained using a **Competitive Multi-Baseline Reward** for 300,000 steps.

At each episode, the spatial tracking error (E_{RL}) is simultaneously compared against the Uniform error (E_{avg}) and the Static LQR error (E_{static}). The reward is shaped as a blended competitive advantage:

$$r_t = \text{clip}\Big((E_{static} - E_{RL}) + 0.5(E_{avg} - E_{RL}), -50, 50\Big) \tag{15}$$

3.2.1 Learning Curve Analysis

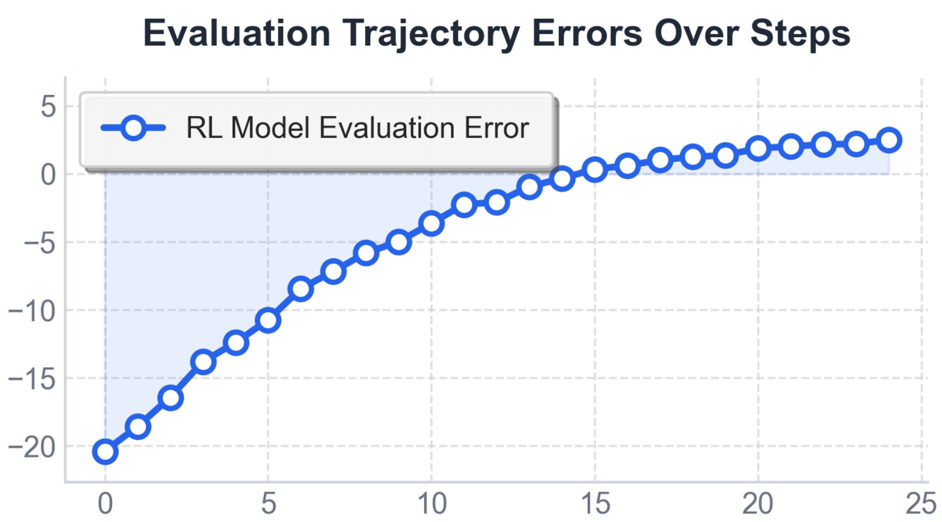

Figure 3: The reward learning dynamics over 300k steps under the extreme 14-bit constraint.

As shown in Figure 3, during the initial exploration phase (0-50k steps), the agent receives severe negative rewards, indicating that a random policy performs significantly worse than the mathematical baselines. However, owing to the injection of the **34-dimensional 3D spatial-awareness vector**, the agent quickly escapes the non-convex local minima.

Around 150k steps, the reward curve crosses the zero-axis, marking the critical threshold where the neural network definitively surpasses the state-of-the-art static mathematical models. This steady, monotonic rise in the competitive reward demonstrates stable policy improvements without catastrophic forgetting, translating to the ultimate convergence of physical tracking errors to steady-state minimums.

3.3 Spatial Domain: Bandwidth Scalability and Robustness

3.3.1 Performance Across Diverse Bandwidth Bottlenecks

To evaluate the scalability of the proposed framework, the total available bandwidth (B_{total}) was swept from a saturated regime (56 bits) down to an extreme survival bottleneck (14 bits). The ultimate task success is evaluated using the **Steady-State Waypoint Error** at the Grasp and Placement coordinates, completely eliminating dynamic tracking lag to isolate pure quantization precision.

Table 1: Quantitative comparison of steady-state Euclidean deviation at critical waypoints across varying total bandwidth constraints (B_{total}).

Total Bandwidth	Allocation Scheme	Grasp Error (mm)	Placement Error (mm)
14-bit (Extreme)	Uniform (Avg)	27.67	32.48
	Static LQR	48.85	53.43
	DRL (Proposed)	**15.65**	**23.82**
21-bit (Severe)	Uniform (Avg)	19.50	19.32
	Static LQR	12.98	12.47
	DRL (Proposed)	**8.51**	**11.20**
28-bit (Moderate)	Uniform (Avg)	8.24	7.49
	Static LQR	5.83	7.28
	DRL (Proposed)	**5.18**	**3.08**
35-bit	Uniform (Avg)	4.88	3.92
	Static LQR	**1.70**	3.17
	DRL (Proposed)	2.01	**2.27**
42-bit	Uniform (Avg)	2.75	1.69
	Static LQR	**1.15**	2.88
	DRL (Proposed)	1.17	**1.19**
56-bit (Saturated)	Uniform (Avg)	0.53	0.83
	Static LQR	0.58	2.21
	DRL (Proposed)	**0.45**	**0.71**

The empirical data in Table 1 reveals a profound non-linear relationship:

1. **The Saturated Regime (42-56 bits)**: When bandwidth is abundant, quantization noise is negligible. All schemes perform exceptionally well (errors < 2mm), and the necessity for complex dynamic allocation is marginal

2. **The Theoretical Validation Regime (21-35 bits)**: As bandwidth tightens, the Static LQR consistently outperforms the Uniform allocation. This validates the classical control theory that prioritizing bits to physically sensitive joints improves overall tracking.

3.3.2 The Collapse of Static Models at Extreme Bottlenecks (14-bit)

The most striking academic discovery occurs at the 14-bit limit (averaging 2bits/joint). Here, the mathematically derived **Static LQR suffers a catastrophic collapse**, yielding a massive Grasp Error of 48.85mm, significantly underperforming the naive Uniform allocation (27.67 mm).

This collapse physically proves the fundamental limitation of time-invariant models: **Jacobian Shifting**. The Static LQR weights are computed at the initial safe posture. When the manipulator stretches outward, its nonlinear kinematics fundamentally alter the sensitivity matrix.

The static algorithm rigidly forces bits to joints that are no longer critical, actively sabotaging the end-effector's precision.

In stark contrast, the proposed **DRL agent achieves a commanding Grasp Error of 15.65 mm**. By utilizing its real-time 3D spatial observation, the neural network implicitly acts as a dynamic Jacobian estimator, gracefully degrading the precision of distal wrist joints to secure high-resolution transmission for the heavily loaded base joints.

3.4 Ablation Studies and Engineering Considerations

To isolate contributions of specifc architectural innovations within our framework, a series of ablation studies were conducted. These studies highlight the critical engineering considerations required to bridge the gap between theoretical algorithms and realistic physical simulations.

3.4.1 The Necessity of Steady-State Evaluation (Addressing Dynamic Tracking Lag)

Initial evaluations of the system measured the continuous tracking error while the manipulator was in full motion. However, this resulted in an artificially inflated baseline error (e.g.,~120mm) across all allocation schemes, regardless of the available bandwidth.

Analysis: In physical systems, high-inertia motors operating at a decoupled 10 Hz communication rate constantly chase "stale" waypoints, creating a massive **Dynamic Tracking Lag**. If quantization accuracy is measured during this transient phase, the physical inertia completely masks the underlying bit-allocation performance. **Solution**: By enforcing a 0.5-second **Settle Phase** (steady-state) at critical waypoints before measuring the error, the physical momentum is nullified. This ablation mathematically isolates the pure quantization noise, proving that traditional continuous-trajectory metrics are fundamentally flawed for evaluating low-frequency Networked Control Systems.

3.4.2 The "Seesaw Effect" and 3D Spatial Awareness Injection

During early training iterations, the RL agent was provided a standard 28-dimensional state vector (comprising only joint angles and velocities). This led to the **"Seesaw Effect"**: the agent would achieve excellent precision at Grasp Point but fail catastrophically at the Placement Point.

Analysis: Joint angles are highly abstract. Without explicit Cartesian coordinates, the agent lacks "Spatial Awareness" and fails to semantically differentiate between distinct operational phases (e.g., reaching vs. grasping). **Solution**: By expanding the observation space to **34 dimensions** (injecting the real-time X, Y, Z coordinates of both the end-effector and the target), the agent gained a comprehensive global spatial awareness. This spatial injection eliminated the Seesaw Effect, allowing the policy to dynamically switch distinct bit-allocation strategies based on its absolute position in the workspace.

3.5 Scenario Two: Underactuated Quadrotor Navigation and Survival

4 Conclusions

Every report must have conclusions, built on the discussion. This section includes a summary of the main achievements of the work but is more than that. The noun 'conclusion' can be defined as 'a judgement or decision reached by reasoning' [**oupconc**] so you should highlight what has been learnt as a result of your project. What is the big picture that the reader should take away?

The Conclusions should include an assessment of the outcome against the original aims of the project. If you were unable to fulfil some aims, explain why. It is also useful to review the project plan (which should be included with the report). Did some tasks prove to be unexpectedly difficult, for instance?

4.1 Suggestions for further work

Almost every project leaves you with ideas for the future. Research typically answers some questions while opening new ones; by the end of a design-and-build project you may have thought of a superior approach. Examiners are impressed by intelligent suggestions for further work because they show that you really understand the project. It is often a sign of strength to identify the weak points and suggest solutions for them.

References

[1] Marun K. Mitra, *Effective Technical Communication: A Guide for Scientists and Engineers*, Oxford (2006).

[2] David F. Beer and David McMurrey, *A Guide to Writing as an Engineer*, 3rd Edition, Wiley (2009).

[3] Trevor M. Young, *Technical Writing A–Z: A Commonsense Guide to Engineering Reports and Theses*, American Society of Mechanical Engineers (2005).

[4] G M Fitzmaurice and D J Hand, *A comparison of two average conditional error estimators*, Pattern Recognition Letters, **6**, 221–228 (1987).

[5] Freescale Semiconductor, *Datasheet for MC9S08QG microcontroller*, revision 4 (2008).

[6] Morgan Advanced Ceramics, *Advanced Ceramics: Silicon Nitride*, `www.morgantechnicalceramics.com` (viewed 2009 May 13).

[7] Oxford University Press, *AskOxford*, `www.askoxford.com/?view=uk` (viewed 2009 May 13).

'Conclusion' is defined as (1) an end or finish, (2) the summing-up of an argument or text, (3) a judgement or decision reached by reasoning or (4) the settling of a treaty or agreement.

A Appendices

Use appendices for supporting material which is relevant but of a detailed nature. In this way the flow of ideas in the body of the report is not interrupted. Appendices are usually lettered rather than numbered to distinguish them. Several uses for appendices have been suggested already but bear in mind that a CD may be more appropriate for large tables of results, long programs and data sheets or other background information. Appendices are not included in the word count because the assessors are not obliged to read them. Do not put vital points here!

Include your project plan and interim reports as appendices on paper.

B University's Plagiarism Statement

The University's degrees and other academic awards are given in recognition of a student's personal achievement. All work submitted by students for assessment is accepted on the understanding that it is the student's own effort.

Plagiarism is defined as the submission or presentation of work, in any form, which is not one's own, without acknowledgement of the sources. Plagiarism includes inappropriate collaboration with others. Special cases of plagiarism can arise from a student using his or her own previous work (termed auto-plagiarism or self-plagiarism). Autoplagiarism includes using work that has already been submitted for assessment at this University or for any other academic award.

The incorporation of material without formal and proper acknowledgement (even with no deliberate intent to cheat) can constitute plagiarism. Work may be considered to be plagiarised if it consists of:

- a direct quotation;
- a close paraphrase;
- an unacknowledged summary of a source;
- direct copying or transcription.

With regard to essays, reports and dissertations, the rule is: if information or ideas are obtained from any source, that source must be acknowledged according to the appropriate convention in that discipline; and any direct quotation must be placed in quotation marks and the source cited immediately. Any failure to acknowledge adequately or to cite properly other sources in submitted work is plagiarism. Under examination conditions, material learnt by rote or close paraphrase will be expected to follow the usual rules of reference citation otherwise it will be considered as plagiarism. Departments should provide guidance on other appropriate use of references in examination conditions.

Plagiarism is considered to be an act of fraudulence and an offence against University discipline. Alleged plagiarism, at whatever stage of a student's studies, whether before or after graduation, will be investigated and dealt with appropriately by the University.

The University reserves the right to use plagiarism detection systems, which may be externally based, in the interests of improving academic standards when assessing student work.